

Lecture 16: PyTorch Crash Course

```
In [1]: import torch
import numpy as np
```

Tensor Initialization and Tensor Properties

See the full API at <https://pytorch.org/docs/stable/tensors.html>

Initializing a 1D tensor from a list

```
In [2]: x=torch.tensor([1,2,3,4])
print(x)

tensor([1, 2, 3, 4])
```

Getting the shape of a tensor

```
In [3]: x.shape

Out[3]: torch.Size([4])
```

Reshaping a tensor

```
In [4]: x.reshape(4,1)

Out[4]: tensor([[1],
               [2],
               [3],
               [4]])
```

```
In [5]: x.reshape(1,4)

Out[5]: tensor([[1, 2, 3, 4]])
```

```
In [6]: x.reshape(1,4).shape

Out[6]: torch.Size([1, 4])
```

Getting the type of a tensor

```
In [7]: type(x)

Out[7]: torch.Tensor
```

```
In [8]: x.dtype

Out[8]: torch.int64
```

Setting the type of a tensor

```
In [9]: x=torch.tensor([1,2,3,4]).type(torch.float)
x
```

```
Out[9]: tensor([1., 2., 3., 4.])
```

```
In [10]: x.dtype
```

```
Out[10]: torch.float32
```

```
In [11]: x=torch.FloatTensor([1,2,3,4])
x
```

```
Out[11]: tensor([1., 2., 3., 4.])
```

```
In [12]: x.dtype
```

```
Out[12]: torch.float32
```

Initializing 2D Tensors

```
In [13]: y = torch.FloatTensor([[1,2,3,4],[5,6,7,8]])
y
```

```
Out[13]: tensor([[1., 2., 3., 4.],
                [5., 6., 7., 8.]])
```

Built in Tensor Constructors

```
In [14]: torch.zeros(5)
```

```
Out[14]: tensor([0., 0., 0., 0., 0.])
```

```
In [15]: torch.ones(5)
```

```
Out[15]: tensor([1., 1., 1., 1., 1.])
```

```
In [16]: torch.rand(5)
```

```
Out[16]: tensor([0.0511, 0.8955, 0.5597, 0.6825, 0.8168])
```

```
In [17]: torch.eye(5)
```

```
Out[17]: tensor([[1., 0., 0., 0., 0.],
                [0., 1., 0., 0., 0.],
                [0., 0., 1., 0., 0.],
                [0., 0., 0., 1., 0.],
                [0., 0., 0., 0., 1.]])
```

Constructing tensors for Numpy arrays

```
In [18]: torch.FloatTensor(np.random.randn(5,3))
```

```
Out[18]: tensor([[ -1.0754, -1.0123, -0.7209],
                [-0.6453, -1.7376, -1.3521],
                [ 1.1395,  0.1085,  0.8397],
                [ 0.5225, -2.1468, -1.3777],
                [ 0.1642,  0.8596,  0.2979]])
```

Linear Algebra

Create data

```
In [19]: x=torch.FloatTensor([[1,2,3,4]])
        z=torch.FloatTensor([[1,2,3,4],[5,6,7,8]])
        print("x=",x)
        print("z=",z)

x= tensor([[1., 2., 3., 4.]])
z= tensor([[1., 2., 3., 4.],
          [5., 6., 7., 8.]])
```

Multiplication by a scalar

```
In [20]: 5*x

Out[20]: tensor([[ 5., 10., 15., 20.]])
```

Addition of a scalar

```
In [21]: x+5

Out[21]: tensor([[6., 7., 8., 9.]])
```

Addition of two tensors of the same shape

```
In [22]: x+x

Out[22]: tensor([[2., 4., 6., 8.]])
```

Transpose of a tensor

```
In [23]: x.T

Out[23]: tensor([[1.],
                [2.],
                [3.],
                [4.]])
```

Matrix Multiplication

```
In [24]: x@x.T

Out[24]: tensor([[30.]])
```

```
In [25]: x.T@x
```

```
Out[25]: tensor([[ 1.,  2.,  3.,  4.],
                 [ 2.,  4.,  6.,  8.],
                 [ 3.,  6.,  9., 12.],
                 [ 4.,  8., 12., 16.]])
```

```
In [26]: x@z.T
```

```
Out[26]: tensor([[30., 70.]])
```

Elementwise multiplication (Hadamard product)

```
In [27]: x*x
```

```
Out[27]: tensor([[ 1.,  4.,  9., 16.]])
```

Broadcasting

Create data

```
In [28]: x=torch.FloatTensor([[1,2,3,4]])
y=torch.FloatTensor([[-1,-2]]).T
z=torch.FloatTensor([[1,2,3,4],[5,6,7,8]])
print("x=",x,"\n")
print("y=",y,"\n")
print("z=",z,"\n")
```

```
x= tensor([[1., 2., 3., 4.]])
```

```
y= tensor([[ -1.],
           [-2.]])
```

```
z= tensor([[1., 2., 3., 4.],
           [5., 6., 7., 8.]])
```

Addition with broadcasting

```
In [29]: print("x=",x,"\n")
print("y=",y,"\n")
a = x+y
print("x+y=",a)
```

```
x= tensor([[1., 2., 3., 4.]])
```

```
y= tensor([[ -1.],
           [-2.]])
```

```
x+y= tensor([[ 0.,  1.,  2.,  3.],
             [-1.,  0.,  1.,  2.]])
```

```
In [30]: print("x=",x,"\n")
print("z=",z,"\n")
a = x+z
print("x+z=",a)

x= tensor([[1., 2., 3., 4.]])

z= tensor([[1., 2., 3., 4.],
          [5., 6., 7., 8.]])

x+z= tensor([[ 2., 4., 6., 8.],
            [ 6., 8., 10., 12.]])
```

```
In [31]: print("y=",y,"\n")
print("z=",z,"\n")
a = y+z
print("y+z=",a)

y= tensor([[ -1.],
          [-2.]])

z= tensor([[1., 2., 3., 4.],
          [5., 6., 7., 8.]])

y+z= tensor([[0., 1., 2., 3.],
            [3., 4., 5., 6.]])
```

Elementwise multiplication with broadcasting

```
In [32]: x*y

Out[32]: tensor([[ -1., -2., -3., -4.],
               [-2., -4., -6., -8.]])

In [33]: x*z

Out[33]: tensor([[ 1., 4., 9., 16.],
               [ 5., 12., 21., 32.]])

In [34]: y*z

Out[34]: tensor([[ -1., -2., -3., -4.],
               [-10., -12., -14., -16.]])
```

Elementwise Functions

```
In [35]: x=torch.FloatTensor([[1,2,-3,-4]])
x

Out[35]: tensor([[ 1., 2., -3., -4.]])

In [36]: torch.pow(x,2)

Out[36]: tensor([[ 1., 4., 9., 16.]])

In [37]: x**2
```

```
Out[37]: tensor([[ 1.,  4.,  9., 16.]])
```

```
In [38]: torch.exp(x)
```

```
Out[38]: tensor([[2.7183, 7.3891, 0.0498, 0.0183]])
```

```
In [39]: torch.log(x)
```

```
Out[39]: tensor([[0.0000, 0.6931,    nan,    nan]])
```

```
In [40]: torch.sin(x)
```

```
Out[40]: tensor([[ 0.8415,  0.9093, -0.1411,  0.7568]])
```

```
In [41]: torch.abs(x)
```

```
Out[41]: tensor([[1., 2., 3., 4.]])
```

Matrix Functions

```
In [42]: z = torch.FloatTensor([[2,1],[1,3]])  
z
```

```
Out[42]: tensor([[2., 1.],  
                [1., 3.]])
```

```
In [43]: torch.det(z)
```

```
Out[43]: tensor(5.)
```

```
In [44]: torch.inverse(z)
```

```
Out[44]: tensor([[ 0.6000, -0.2000],  
                [-0.2000,  0.4000]])
```

```
In [45]: torch.trace(z)
```

```
Out[45]: tensor(5.)
```

Aggregation Functions

Create data

```
In [46]: x=torch.FloatTensor([[1,2,3,4]])  
z=torch.FloatTensor([[1,2,3,4],[5,6,7,8]])  
print("x=",x,"\n")  
print("z=",z)
```

```
x= tensor([[1., 2., 3., 4.]])
```

```
z= tensor([[1., 2., 3., 4.],  
          [5., 6., 7., 8.]])
```

Summing over all tensor elements

```
In [47]: x.sum()
```

```
Out[47]: tensor(10.)
```

```
In [48]: z.sum()
```

```
Out[48]: tensor(36.)
```

Summing over a specific axis

```
In [49]: z.sum(axis=0)
```

```
Out[49]: tensor([ 6.,  8., 10., 12.])
```

```
In [50]: z.sum(axis=1)
```

```
Out[50]: tensor([10., 26.])
```

```
In [51]: z.sum(axis=1, keepdims=True)
```

```
Out[51]: tensor([[10.],  
                [26.]])
```

Mean and Product

```
In [52]: x.mean()
```

```
Out[52]: tensor(2.5000)
```

```
In [53]: x.prod()
```

```
Out[53]: tensor(24.)
```

Min and Max

```
In [54]: x.min()
```

```
Out[54]: tensor(1.)
```

```
In [55]: x.max()
```

```
Out[55]: tensor(4.)
```

```
In [56]: z.min()
```

```
Out[56]: tensor(1.)
```

```
In [57]: v,ind = z.min(axis=0)
print(v)
print(ind)

tensor([1., 2., 3., 4.])
tensor([0, 0, 0, 0])
```

```
In [58]: v,ind = z.min(axis=1)
print(v)
print(ind)

tensor([1., 5.])
tensor([0, 0])
```

Automatic Differentiation

Initializing tensors

```
In [59]: x=torch.FloatTensor([[1,2,3,4],[-1,-2,-3,-4]])
y=torch.tensor([1.0],[0.0])
print("x=",x,"\n")
print("y=",y,"\n")

x= tensor([[ 1.,  2.,  3.,  4.],
            [-1., -2., -3., -4.]])

y= tensor([[1.],
            [0.]])
```

Initializing tensors with gradient tracking

```
In [60]: w=torch.tensor([[-1.0],[-2.0],[2.0],[1.0]],requires_grad=True)
b=torch.tensor(-1.0,requires_grad=True)
print("w=",w,"\n")
print("b=",b,"\n")

w= tensor([[-1.],
            [-2.],
            [ 2.],
            [ 1.]], requires_grad=True)

b= tensor(-1., requires_grad=True)
```

Checking if gradient tracking is enabled

```
In [61]: x.requires_grad
```

```
Out[61]: False
```

```
In [62]: w.requires_grad
```

```
Out[62]: True
```



```
In [63]: b.requires_grad
```

```
Out[63]: True
```

Performing Operations

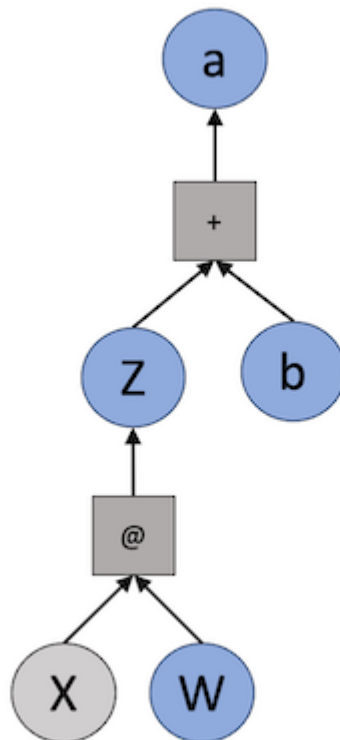
```
In [64]: a=x[0,:].@w +b  
a
```

```
Out[64]: tensor([4.], grad_fn=<AddBackward0>)
```

```
In [65]: a.requires_grad
```

```
Out[65]: True
```

Computation Graph



Getting gradient values

```
In [66]: a.backward()  
print("grad_w:",w.grad,"\n")  
print("grad_b:",b.grad)
```

```
grad_w: tensor([[1.],
               [2.],
               [3.],
               [4.]])
```

```
grad_b: tensor(1.)
```

Gradient accumulation and clearing

```
In [67]: a=x[0,:].@w +b
         a.backward()
         print("grad_w:",w.grad,"grad_b:",b.grad)
```

```
grad_w: tensor([[2.],
               [4.],
               [6.],
               [8.]]) grad_b: tensor(2.)
```

```
In [68]: w.grad=None
         b.grad=None
         a=x[0,:].@w +b
         a.backward()
         print("grad_w:",w.grad,"grad_b:",b.grad)
```

```
grad_w: tensor([[1.],
               [2.],
               [3.],
               [4.]]) grad_b: tensor(1.)
```

Hiding operations from gradient tracking

```
In [69]: w.grad=None
         b.grad=None
         with torch.no_grad():
             a=x[0,:].@w +b
         a
```

```
Out[69]: tensor([4.])
```

Example: OLS gradient

```
In [70]: def mse(y,yhat):
         return torch.mean((y-yhat)**2)

         w.grad=None
         b.grad=None

         yhat = x@w+b
         loss = mse(y,yhat)
         loss.backward()
         print("grad_w:",w.grad,"\n")
         print("grad_b:",b.grad)
```

```
grad_w: tensor([[ 9.],
                [18.],
                [27.],
                [36.]])
```

```
grad_b: tensor(-3.)
```

Neural Network Modules

See the full API at <https://pytorch.org/docs/stable/nn.html>

Generate data

```
In [71]: X = torch.randn(1000,2)
Y = -3*X[:,[0]] + 2*X[:,[1]]**2 + 3
Y[:5,:]
```

```
Out[71]: tensor([[ 1.3281],
                [ 4.5852],
                [ 1.1751],
                [ 5.3190],
                [12.0424]])
```

Defining a linear layer from scratch

```
In [72]: import torch.nn as nn

class linear(nn.Module):
    def __init__(self,d,k):
        super(linear, self).__init__()
        self.w = nn.Parameter(torch.rand(d,k))
        self.b = nn.Parameter(torch.rand(1,k))

    def forward(self,x):
        return x@self.w + self.b
```

Instantiating the layer and computing output

```
In [73]: model = linear(2,1)
model.forward(X[:10,:])
```

```
Out[73]: tensor([[ 0.7800],
                [ 0.3021],
                [ 0.6498],
                [ 0.1447],
                [-0.1138],
                [-0.2512],
                [ 0.1483],
                [ 0.1373],
                [ 0.2449],
                [ 0.6503]], grad_fn=<AddBackward0>)
```

Computing a loss on the output

```
In [74]: Yhat = model.forward(X)
         loss = mse(Yhat,Y)
         loss
```

```
Out[74]: tensor(39.8705, grad_fn=<MeanBackward0>)
```

Manually setting parameters

```
In [75]: model.w.data = torch.zeros(2,1)
         model.b.data = torch.zeros(1,1)
         print("w:",model.w)
         print("b:",model.b)
```

```
w: Parameter containing:
tensor([[0.],
        [0.]], requires_grad=True)
b: Parameter containing:
tensor([[0.]], requires_grad=True)
```

```
In [76]: model.forward(X[:10,:])
```

```
Out[76]: tensor([[0.],
                 [0.],
                 [0.],
                 [0.],
                 [0.],
                 [0.],
                 [0.],
                 [0.],
                 [0.],
                 [0.]], grad_fn=<AddBackward0>)
```

Defining a relu MLP using built-in layers

```
In [77]: class relu_mlp(nn.Module):
         def __init__(self,d,k):
             super(relu_mlp, self).__init__()
             self.l1 = nn.Linear(d,k)
             self.relu = nn.ReLU()
             self.out = nn.Linear(k,1)

         def forward(self,x):
             return self.out(self.relu(self.l1(x)))
```

Inspecting the model

```
In [78]: model = relu_mlp(2,3)
         print("l1:", model.l1)
         print("relu:",model.relu)
         print("out:",model.out)
```

```
l1: Linear(in_features=2, out_features=3, bias=True)
relu: ReLU()
out: Linear(in_features=3, out_features=1, bias=True)
```

Inspecting the parameters

```
In [79]: print("l1 weight", model.l1.weight, "\n")
         print("l1 bias", model.l1.bias)

l1 weight Parameter containing:
tensor([[ 0.0766, -0.4320],
        [-0.6615, -0.0574],
        [-0.4294,  0.2210]], requires_grad=True)

l1 bias Parameter containing:
tensor([-0.4012, -0.5149, -0.4955], requires_grad=True)
```

Computing output

```
In [80]: model.forward(X[:10,:])

Out[80]: tensor([[0.4863],
                [0.4863],
                [0.4863],
                [0.4863],
                [0.4823],
                [0.6364],
                [0.4834],
                [0.6093],
                [0.4862],
                [0.4862]], grad_fn=<AddmmBackward0>)
```

Computing a loss on the output

```
In [81]: Yhat = model.forward(X)
         loss = mse(Yhat, Y)
         loss

Out[81]: tensor(36.7757, grad_fn=<MeanBackward0>)
```

Manipulating the parameters built-in layers

```
In [82]: with torch.no_grad():
         model.l1.weight.data = torch.zeros_like(model.l1.weight.data)
         model.l1.bias.data   = torch.zeros_like(model.l1.bias.data)
         model.out.weight.data = torch.zeros_like(model.out.weight.data)
         model.out.bias.data   = torch.zeros_like(model.out.bias.data)

         print("l1 weight:", model.l1.weight, "\n")
         print("l1 bias:", model.l1.bias, "\n")
         print("out weight:", model.out.weight, "\n")
         print("out bias:", model.out.bias, "\n")
```

```

l1 weight: Parameter containing:
tensor([[0., 0.],
        [0., 0.],
        [0., 0.]], requires_grad=True)

l1 bias: Parameter containing:
tensor([0., 0., 0.], requires_grad=True)

out weight: Parameter containing:
tensor([[0., 0., 0.]], requires_grad=True)

out bias: Parameter containing:
tensor([0.], requires_grad=True)

```

```
In [83]: model.forward(X[:10,:])
```

```

Out[83]: tensor([0.,
                0.,
                0.,
                0.,
                0.,
                0.,
                0.,
                0.,
                0.,
                0.], grad_fn=<AddmmBackward0>)

```

Optimization

Defining a basic fit function

```

In [84]: def fit(model, lr, max_iter):
    optimizer = torch.optim.SGD(model.parameters(), lr=lr)
    losses=[]
    for i in range(max_iter):
        Yhat = model.forward(X)
        loss = mse(Yhat,Y)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        losses.append(loss.detach().numpy().item())
        if(i%10==0):
            print("%d %.2f"%(i, losses[-1]))

    return(losses)

```

Instantiating models

```
In [85]: models = {"Linear":linear(2,1),"RELU MLP":relu_mlp(2,20)}
```

Fitting models

```
In [86]: losses = {}
        for m in models:
            print("Learning model: %s"%m)
            losses[m] = fit(models[m], 0.05, 100)
            print()
```

Learning model: Linear

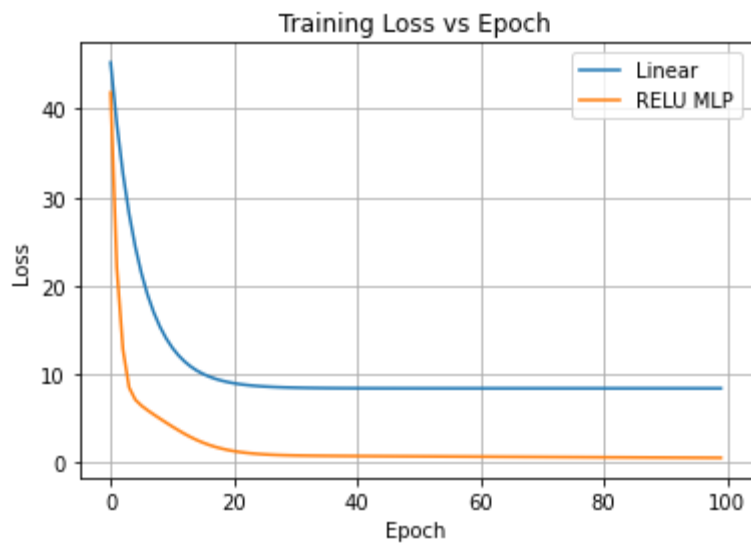
0 45.25
10 13.00
20 8.98
30 8.48
40 8.41
50 8.41
60 8.40
70 8.40
80 8.40
90 8.40

Learning model: RELU MLP

0 41.85
10 4.09
20 1.30
30 0.81
40 0.75
50 0.71
60 0.68
70 0.64
80 0.61
90 0.57

Plotting losses

```
In [87]: import matplotlib.pyplot as plt
        plt.figure()
        for m in models:
            plt.plot(losses[m])
        plt.xlabel("Epoch")
        plt.ylabel("Loss")
        plt.legend(list(models.keys()))
        plt.title("Training Loss vs Epoch")
        plt.grid(True)
```



Computing error of fit models with eval and no_grad

```
In [88]: Xte = torch.randn(1000,2)
Yte = -3*Xte[:,[0]] + 2*Xte[:,[1]]**2 + 3

te_err={}
for m in models:
    models[m].eval()
    with torch.no_grad():
        te_err[m]=mse(Yte,models[m].forward(Xte))
        print("%s test error: %.2f"%(m,te_err[m]))
```

Linear test error: 9.36
RELU MLP test error: 0.73

In []: